

Discrete Probability Distribution of Penguin Weights

I created a discrete probability distribution (Table 1) and density curve with superimposed normal distribution (Figure 1) for the weights of penguins in the PalmerPenguins dataset. The mean weight in the data is approximately equal to 4202 grams, with $P(4202)$ approximately equal to 3.5%. The probabilities and weights are not normally distributed. The Q-Q plot confirms this in Figure 2.

A tibble: 94 × 2			
weights	probabilities		
<dbl>	<dbl>		
2700	0.0029	3650	0.0175
2850	0.0058	3675	0.0058
2900	0.0117	3700	0.0322
2925	0.0029	:	:
2975	0.0029	4775	0.0029
3000	0.0058	4800	0.0088
3050	0.0117	4850	0.0117
3075	0.0029	4875	0.0088
3100	0.0029	4900	0.0058
3150	0.0117	4925	0.0058
3175	0.0058	4950	0.0058
3200	0.0146	4975	0.0029
3250	0.0146	5000	0.0175
3275	0.0029	5050	0.0088
3300	0.0175	5100	0.0088
3325	0.0146	5150	0.0058
3350	0.0146	5200	0.0117
3400	0.0234	5250	0.0088
3425	0.0058	5300	0.0117
3450	0.0234	5350	0.0088
3475	0.0088	5400	0.0146
3500	0.0205	5450	0.0029
3525	0.0058	5500	0.0146
3550	0.0263	5550	0.0175
3575	0.0029	5600	0.0058
3600	0.0205	5650	0.0088
3625	0.0029	5700	0.0146
		5750	0.0029
		5800	0.0058
		5850	0.0088
		5950	0.0058
		6000	0.0058
		6050	0.0029
		6300	0.0029

Table 1




Figure 1




Figure 2

Binomial Probability Distribution and Histogram

I made a binomial probability distribution (and histogram) for the species Adelie in the PalmerPenguins dataset. See Figure 3. This shows the probabilities of the first 1-30 penguins selected being the species Adelie. Increasing the number to 30 more closely approximates the normal distribution than 10 (Figure 4).




Figure 3

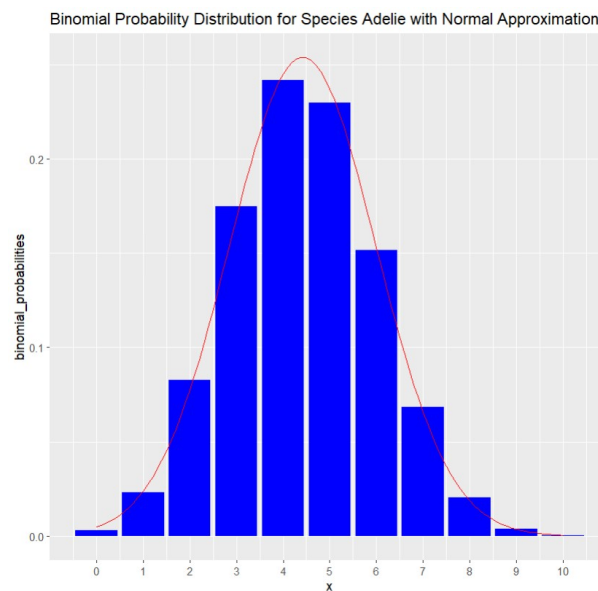


Figure 4

Visualization:

- The histogram visualizes the data from the table, with the x-axis representing the number of successes and the y-axis representing the probability of that number of successes.
- Each bar's height corresponds to the probability of that number of successes occurring. For instance, the tallest bar at $x=13$ corresponds to the highest probability (about 0.14) among the probabilities listed in the table, meaning that observing exactly 13 successes is the most probable outcome within this dataset.
- This visualization is particularly useful in showing the distribution of probabilities across different numbers of successes. It provides a quick way to see which outcomes are most and least likely and how the probabilities are distributed (e.g., if they are skewed towards lower or higher numbers of successes).

I made no changes to my dataset this week.